

<!-- v1.8 -->

PRIVATE AUDIO BRIEFING / AI Roadmap for Scott Paul - Phi Health / Frequency / Prepared by Franchise Ki

Generated: 2026-06-18T00:00:00Z

Open the audio with EXACTLY these words: "Hi Scott, welcome. Scott, this walkthrough was built for you and Phi Health / Frequency from the public Phi materials."

DIRECT ADDRESS RULE: The audio must open with: Hi Scott, welcome. Scott, this walkthrough was built for you and Phi Health / Frequency from the public Phi materials.

If any opening phrase gets shortened by the audio host, the first 15 seconds must still contain the words "Scott", "welcome", "you", and "Phi Health / Frequency".

Do not begin with "Frequency from the public materials", "thanks for having me", "right now you're sitting on", or any phrase that omits Scott's name.

Audience: Scott Paul, Phi Health / Frequency.

Speak TO Scott as "you" and "your" for the entire episode.

STRICT BANS: do not use third-person owner, company, or business labels for Scott or Phi Health. Say "you", "your brand", "your podcast", "your Phi ecosystem", "your workflow", and "your audience".

Never mention internal audit, production gates, background files, customer send, Bennett preview, or where this briefing came from.

NotebookLM length: SHORT. Target 6 to 8 minutes. Do not cut off mid-thought. End cleanly after the next-step summary.

CTA: complete the qualifier only; no calendar booking language.

SOURCE FACTS TO CARRY FORWARD

UPDATED PHI ECOSYSTEM POSITIONING — USE THIS OVER OLDER PODCAST-ONLY FRAMING

- This is not only a podcast roadmap. Talk to Scott as someone building an AI-native health ecosystem.
- Confirmed public surfaces: Phi website, Frequency, Episode 1, Phi Stack, Phi iOS App / AI Health Journal, shop, research/blog, support/contact, and the Episode 1 Stack giveaway.
- Frame Frequency as the expert-input layer. Each episode should feed structured knowledge into app education, Stack-safe learning, clips, FAQs, support macros, and listener routing.
- The strongest point is that Scott already understands AI, so generic chatbot advice is weak. The useful angle is connecting existing Phi assets into a specific operating loop.
- Private metrics are unknown. Do not say lead volume, revenue, app usage, support volume, close rate, or conversion numbers as facts. Treat any calculator numbers as editable scenarios only.
- The target audience evidence is public/inferred: healthspan-minded listeners, Stack customers, AI health journal users, people interested in body/mind, recovery, sleep, energy, and expert-led health education. Action-sports and athlete language should be tied to Dr. Kramer's Episode 1 context, not claimed as the whole audience.
- The hosts should say this directly to Scott: your fastest win is the episode-to-app-to-Stack loop: prep the guest, map the transcript, review claims, create clips, route giveaway responses, and hand off app/Stack questions safely.
- Never use source-document labels, report labels, case-study labels, or intake-style framing. Say "from the public Phi materials" or simply speak directly.

SHORT AUDIO SHAPE

Open directly: "Hi Scott, welcome. Scott, this walkthrough was built for you and Phi Health / Frequency from the public Phi materials."

Then cover only five points:

1. Phi is AI Health Journal + Stack + Frequency, not just a podcast.
2. Frequency can become the knowledge input layer for app education and Stack-safe learning.

3. The first episode has strong raw material: Dr. Kramer, 50 minutes, recovery, PRP, fascia, D3/K2, absorption, sleep/aging, and giveaway response.
 4. The six-agent system: guest prep, episode knowledge graph, claims review, clips, listener/giveaway intake, app/Stack learning path.
 5. Next step: if this resonates, open the written Blueprint and fill out the qualifier so the real numbers and internal workflow can replace assumptions.
- End cleanly. Do not trail off. Do not mention audio length.

UPDATED SOURCE SNAPSHOT GENERATED 2026-06-18T15:53:59.153547+00:00

Frequency page:

- Public URL: <https://www.phi.health/podcasts/frequency>
- Page title: Frequency.
- Positioning: A Phi Podcast.
- Core promise shown on the page: learn from the best in the world about body and mind.
- Core question shown on the page: how do we find more healthy years, lived at a higher frequency, doing the things we love, with the people we love?
- The first listed episode is "Lessons on Healing from 35 Years of Orthopedic Surgery #1 Dr. Warren G. Kramer III, MD".
- The first listed episode is shown as 50 minutes and dated June 9, 2026.

Episode page:

- Public URL:
<https://www.phi.health/podcasts/frequency/1-lessons-on-healing-from-35-years-of-orthopedic-surgery-dr-warren-g-kramer-iii-md-g6mas9w>
- Episode topic: recovery, aging, nutrition, the body's healing capacity, and lessons from decades of orthopedic surgery.
- "What you'll learn" themes shown on the page include: rest and icing can slow recovery; PRP can restart healing; muscles protect bones; vitamin D3 works best with K2 and absorption; movement practices like yoga, Pilates, and Tai Chi are discussed.
- Chapters include the start of the conversation, alternative approaches, fascia, stem cells and PRP, nutrition, supplements, and final reflections.

IMPORTANT HEALTH-CONTENT RULES

This is health media. Do not make personal medical recommendations. Do not say that AI should diagnose, treat, prescribe, dose, cure, prevent disease, or replace a clinician. Keep all language educational, careful, and review-oriented. When discussing the Phi Stack, supplements, D3, K2, PRP, stem cells, fascia, recovery, or aging, frame the workflow as "organize, summarize, review, cite, and route" instead of "advise, recommend, or prescribe."

Use phrases like:

- "educational language"
- "human review"
- "qualified professional"
- "approved source"
- "claim-risk note"
- "citation needed"
- "listener-first takeaway"

Avoid phrases like:

- "the AI tells listeners what to take"
- "the agent recommends treatment"
- "guaranteed recovery"

- "proven cure"
- "medical advice"

12 SECTION AUDIO PACKAGE

1. About Scott and Phi Health / Frequency

You have a clear media premise: Frequency is a Phi podcast about learning from world-class voices on body and mind. The public page asks a useful question: how do people find more healthy years, lived at a higher frequency, doing what they love with people they love? That gives the podcast an emotional center and a practical content lane. It is not just interviews; it is a trusted education channel that can feed the broader Phi ecosystem: the AI Health Journal, Stack, research, shop, support, and community follow-up.

The first episode gives a strong operating example. You feature Dr. Warren G. Kramer, an orthopedic surgeon with decades of experience, in a 50-minute episode about healing, recovery, aging, nutrition, regenerative medicine, fascia, movement, D3, K2, and absorption. The episode is rich enough to support show notes, clips, newsletter sections, app education, guest follow-up, and safe product-adjacent education. It also carries health-claim sensitivity, which means the workflow has to be careful.

2. Current Tool Stack

Your current visible stack includes the Phi website, the Frequency podcast page, episode pages, transcripts, podcast audio, social distribution paths, the Phi AI Health Journal path, a shop and product education path around Phi Stack, and likely a support/contact layer for listeners and customers. The system should sit around those pieces. It should not start with a platform migration.

The right AI layer should attach to:

- episode research and guest prep
- transcript ingestion and knowledge graph tagging
- show note, chapter, app snippet, and FAQ production
- claim and citation review
- clip and social repurposing
- listener question routing
- app and Stack learning path handoff
- giveaway or community engagement workflows

3. Industry Context

Health media is a high-trust category. Listeners are not just browsing entertainment; they are thinking about their bodies, recovery, energy, aging, pain, performance, sleep, supplements, or the quality of their future life. That attention is valuable, and it raises the bar for accuracy.

The best use of AI in this context is not to create clinical answers. It is to reduce production drag while increasing review discipline. AI can help you prepare better questions, structure transcripts, surface the best moments, write careful summaries, flag claims, and route messages quickly. It should help your team be more useful and more careful at the same time.

4. Where Time Is Leaking

The first leak is post-production. A long expert episode contains dozens of usable moments, but extracting them by hand takes time. Show notes, chapters, clips, newsletters, app snippets, SEO summaries, and social captions can all come from the same transcript if the workflow is set up correctly.

The second leak is claims review. A phrase that is fine inside a long expert discussion can become risky when it is separated into a caption, product education note, or email. You need a systematic review pass that flags supplement claims, treatment claims, dosing claims, cure/prevention language, diagnosis-adjacent wording, and sentences that need evidence.

The third leak is listener routing. Someone may ask about the episode, a guest, the giveaway, a product, app access, an order, or a personal health concern. Those should not all go into one generic inbox. The faster each question is routed, the more useful the brand feels.

The fourth leak is content library structure. Every episode can build a searchable knowledge base if it is tagged at the transcript level. Without that, valuable ideas disappear into a feed.

5. Six Recommended Agents

Recommended agents:

1. Guest Research and Interview Prep Agent.
2. Episode Knowledge Graph Agent.
3. Claims and Citation Review Agent.
4. Clip and Social Repurposing Agent.
5. Listener Intake and Giveaway Ops Agent.
6. App and Stack Learning Path Agent.

Each agent should have approval boundaries. None of them should publish, send, or provide individualized health guidance without human review.

6. Agent 1: Guest Research and Interview Prep Agent

Before each recording, this agent prepares a guest dossier, topic map, careful claim notes, and listener-first questions. For the Dr. Kramer episode, the prep structure would have included recovery myths, the role of rest and icing, the conversation around PRP and stem cells, fascia, movement practices, muscle protecting bone, D3, K2, and absorption. It would not invent studies or credentials. It would flag claims that need review.

Outcome for Phi Health / Frequency: each episode starts with sharper questions, fewer missed follow-ups, and a clearer arc for the listener.

Prompt:

You are the Guest Research and Interview Prep Agent for Scott Paul and Phi Health / Frequency. Build a guest dossier from approved public sources and internal notes. Output guest background, strongest topics, five listener-first questions, three contrarian follow-ups, claim-risk notes, and a plain-language episode thesis. Keep the tone curious and rigorous. Do not invent credentials, studies, quotes, or medical facts. Flag any claim that needs human review.

7. Agent 2: Episode Knowledge Graph Agent

After each episode, this agent turns the transcript into chapter notes, show notes, listener takeaways, newsletter blocks, app-friendly snippets, SEO summaries, and clip candidates. The first episode can become a structured education package with themes like recovery conditions, movement, fascia, muscle, regenerative therapy discussion, nutrition, and nutrient absorption.

Outcome for Phi Health / Frequency: a 50-minute expert interview becomes a full content package instead of one audio asset.

Prompt:

You are the Episode Knowledge Graph Agent for Phi Health / Frequency. Given an episode transcript, create chapter notes, a listener takeaway list, a newsletter section, SEO summary, short-form clip candidates, and app-friendly education snippets. Preserve the guest's meaning, avoid exaggeration, and use educational language. Never convert a general discussion into personal medical advice. Mark statements that need citations or clinician review.

8. Agent 3: Claims and Citation Review Agent

This agent reviews show notes, captions, emails, product education, app copy, and support answers before anything becomes public. It marks medical claims, supplement claims, diagnosis language, treatment language, cure/prevention language, dosing statements, and outcome promises. It suggests safer educational wording and flags when a human should review.

Outcome for Phi Health / Frequency: useful education moves faster while the team avoids turning expert conversation into unsupported public claims.

Prompt:

You are the Claims and Citation Review Agent for Phi Health. Review draft podcast notes, captions, emails, app copy, and product education for unsupported medical, supplement, treatment, cure, prevention, diagnosis, dosing, and outcome claims. Return a table with original sentence, risk level, issue type, safer educational rewrite, citation needed yes/no, and human review needed yes/no. Do not approve claims you cannot support.

9. Agent 4: Clip and Social Repurposing Agent

This agent finds the strongest short moments in each episode and prepares hooks, captions, titles, thumbnail phrases, and channel-specific variants. It should not chase sensational promises. It should use curiosity and clarity. For the Dr. Kramer episode, clips could be built around recovery myths, fascia, PRP discussion, how movement fits into healing, muscle and bone protection, and the nuance around D3, K2, and absorption.

Outcome for Phi Health / Frequency: every recording gets more reach without requiring Scott to manually cut and rewrite every asset.

Prompt:

You are the Clip and Social Repurposing Agent for Frequency. Given an episode transcript and timestamps, identify 12 short clip candidates with title, hook, caption, thumbnail phrase, audience, and claim-risk note. Keep each clip faithful to the source. Prefer curiosity and usefulness over sensational claims. Do not write disease-treatment promises, supplement promises, or clinician-like recommendations.

10. Agent 5: Listener Intake and Giveaway Ops Agent

This agent classifies incoming messages. A listener may have a podcast question, a guest request, a giveaway entry, a Phi Stack question, an app support issue, an order question, a partnership inquiry, or a medical advice request. Each message needs a different route.

Outcome for Phi Health / Frequency: listener interest is handled while it is fresh, and sensitive questions go to the right place.

Prompt:

You are the Listener Intake and Giveaway Ops Agent for Phi Health and Frequency. Classify incoming messages into podcast question, guest request, giveaway, Phi Stack/product education, app support, order/support, partnership, or medical advice request. Draft a short response for human approval. Use safe educational language, link to the right approved page when available, and route medical advice requests to a qualified professional instead of answering clinically.

11. Agent 6: App and Stack Learning Path Agent

This agent answers common non-clinical questions from approved FAQ, shop, app, and podcast materials. It can point a listener to an episode, explain where to learn more, or route an order/app issue. It should never diagnose, prescribe, recommend dosing, or say a product treats or prevents disease.

Outcome for Phi Health / Frequency: listeners get help quickly without pushing AI into clinical advice.

Prompt:

You are the App and Stack Learning Path Agent for Phi Health. Answer common non-clinical questions using approved FAQ, shop, podcast, and app materials. Keep responses educational and concise. Never diagnose, prescribe, recommend dosing, or claim a product treats or prevents disease. If the user asks for personal medical guidance, advise them to consult a qualified clinician and route to human support if appropriate.

12. ROI Picture and Next Step

Use placeholder planning numbers only until Scott supplies actual numbers: about 1,200 monthly listener/product/support interactions, \$150 average commerce or engagement value, 8 percent conversion, and about 30 weekly production/support hours. Do not claim these are Scott's real metrics. The ROI message should be framed around capacity and workflow quality: fewer manual hours per episode, more assets per recording, faster listener routing, and stronger review discipline.

Close by inviting Scott to review the written Blueprint and complete the qualifier. The qualifier lives at <https://bennett-maxwell.github.io/fki-preview/qualify.html>. Mention it as the qualifier page; do not turn this into a calendar-booking CTA.

FULL SPOKEN SCRIPT

Hi Scott, welcome. This walkthrough was built for you and Phi Health / Frequency from the public Phi materials.

I am going to keep this direct and practical. Frequency already has a clear premise: learn from the best people in the world about body and mind, and help people find more healthy years lived at a higher frequency, doing the things they love with the people they love. That is a strong media lane because it is not just content. It is trust, education, and a path into the broader Phi ecosystem.

The first episode gives us a useful model. You have Dr. Warren G. Kramer, an orthopedic surgeon with decades of experience, talking through recovery, aging, nutrition, PRP, fascia, muscle, vitamin D3, K2, and the idea that the body often needs the right conditions to heal. It is rich material. It is also exactly the kind of health content that needs a smarter workflow after the recording is done, because the difference between useful education and risky overstatement can be one sentence in a caption, email, or product note.

So the AI system I would build for you is not a chatbot sitting on top of a health brand. It is a production and trust system around Frequency. It helps you prep the guest, extract the strongest listener takeaways, repurpose the episode, review claim language, route listener interest, and connect people to the right Phi path without pretending to give personal medical advice.

The first agent is your Guest Research and Interview Prep Agent. Before each recording, it builds a guest dossier, topic map, careful claim notes, and listener-first questions. For a guest like Dr. Kramer, it would organize the arc before you sit down: recovery myths, why rest and icing can be misunderstood, when regenerative therapies are discussed, why muscle matters for bones, and where nutrition or absorption belongs in the conversation. It would also mark the places where a claim needs extra care. That means you walk into the interview with sharper follow-ups instead of a generic wellness outline.

The second agent is your Episode Knowledge Graph Agent. After recording, it takes the transcript and turns it into structured knowledge your audience and team can actually use: topic tags, chapter summaries, show notes, quote banks, claim-risk notes, app-friendly snippets, Stack-safe education paths, FAQs, and short clip candidates. A 50-minute episode should not live as one audio file and a title. It should become a searchable knowledge layer for Frequency, the Phi app, Stack education, and support. For your first Frequency episode, that means each major theme can become a clean takeaway: recovery is not always about stopping movement, muscle helps protect structure, absorption matters when talking about nutrients, and medical decisions belong with a qualified professional.

The third agent is your Claims and Citation Review Agent. This is the most important guardrail because health content carries a different responsibility than ordinary business content. A line that feels natural in a conversation can become too strong when it is pulled into an Instagram caption, a product education page, or an email. This agent reviews drafts for treatment, diagnosis, supplement, cure, prevention, dosing, and outcome language. It does not approve anything by itself. It flags risk, proposes safer educational wording, asks for citations where needed, and pushes human review when the content gets clinical.

The fourth agent is your Clip and Social Repurposing Agent. It finds the strongest moments in each episode and drafts the hook, caption, title, thumbnail phrase, and platform-specific angle. The key is that it does not chase sensational health promises. It turns the episode into useful curiosity. For example, the Dr. Kramer episode can generate clips around recovery myths, the role of fascia, why muscles matter, what people misunderstand about vitamin D, and how aging changes the way people should think about movement. Each clip comes with a claim-risk note so the review pass happens before anything goes live.

The fifth agent is your Listener Intake and Giveaway Ops Agent. When someone reacts to an episode, enters a giveaway, asks about a guest, asks about the Phi Stack, has an app issue, or sends a question that sounds personal and medical, this agent sorts it. It drafts a safe response for review and routes the message to the right path. A podcast question does not belong in the same lane as an order question. A product education question does not belong in the same lane as a clinical question. The value here is speed and clarity while listener interest is fresh.

The sixth agent is your App and Stack Learning Path Agent. This one connects the podcast audience to the rest of Phi without crossing medical lines. It maps listener interest to approved Phi app prompts, Stack education, FAQs, shop support, or human escalation. When someone asks for personal medical guidance, the agent does not try to be clever. It says that is a clinician question and routes appropriately.

Here is the rollout I would use.

In the first three days, build the repeatable episode workflow. Set up the guest prep template, transcript intake, claims categories, and routing rules. Your first win is not a huge platform rebuild. It is a clean path from guest research to final content.

In week one and week two, process the Dr. Kramer episode through the system. Pull the transcript into the education agent. Generate the show notes, chapters, newsletter section, app snippets, and clip list. Then run every draft through the claims review agent before anything becomes a public caption or product-adjacent copy. That turns your first Frequency episode into a useful content library instead of a single episode page.

In weeks three and four, scale the format. Add guest dossiers for upcoming episodes, build a clip calendar, and connect listener intake to the right support or education queues. Once the workflow is stable, every new episode gets easier to publish, easier to repurpose, and easier to trust.

The practical reason this matters is that your content sits in a high-trust category. People will listen because they care about their body, recovery, aging, energy, and quality of life. That attention is valuable, but it also means the follow-up needs to be careful. You want helpful education, not unsupported claims. You want fast listener responses, not a messy inbox. You want clips that carry the guest's meaning, not clips that exaggerate it.

There is also a compounding content advantage here. Once the first episode is tagged and reviewed correctly, the next episode does not start from zero. The system starts to know your categories: recovery, mobility, nutrition, supplementation, mindset, guest background, listener questions, product education, and support. That means your team can search prior episodes, reuse approved explanations, and compare what one guest said against another without manually digging through transcripts.

This matters for the audience experience. A listener may not remember the exact episode title, but they may remember a question like, "What did the orthopedic surgeon say about icing?" or "Where did the discussion about vitamin D and K2 happen?" A structured transcript layer lets you answer that kind of question quickly. It can point to the right chapter, summarize the point carefully, and avoid turning a general conversation into a personal recommendation.

It also matters for how you protect the brand. When you build in review from the beginning, you can move faster without becoming sloppy. The claims agent does not slow the content down; it gives the rest of the workflow confidence. Your clip agent can draft more clips because every clip has a claim-risk field. Your education agent can draft more summaries because every summary gets checked. Your intake agent can respond faster because it knows when to route instead of answer.

So the first milestone is not "automate everything." The first milestone is to make the Dr. Kramer episode the template. Once that package is clean, you can repeat the same motion for every new guest: prep, record, transcript, education assets, claims pass, clips, listener routing, and support handoff.

Scott, the bottom line is simple. Frequency has the raw material: expert conversations, a clear wellness premise, and a Phi ecosystem people can continue into. The bottleneck is turning each conversation into a repeatable, reviewed, multi-channel system. That is exactly where AI should sit for you. It prepares the episode, structures the transcript, expands the content, checks the claims, routes the listener, and protects your voice.

When you are ready, the next step is to review the written Blueprint and complete the qualifier. That gives us the real operating numbers, the exact publishing stack, and the approval boundaries so the first live agent can be built around your actual process.

UPDATED LEAD PROFILE APPENDIX — CURRENT BLUEPRINT COPY

<!-- v1.8 -->

PRIVATE AUDIO BRIEFING / AI Roadmap for Scott Paul — Phi Health / Frequency / Prepared by Franchise Ki

Generated: 2026-06-18T15:52:49.341264+00:00

Open the audio with EXACTLY these words: "Hi Scott, welcome. This walkthrough was built for you and Phi Health / Frequency, from the public Phi materials."

DIRECT ADDRESS RULE: The audio must open with: Hi Scott, welcome. This walkthrough was built for you and Phi Health / Frequency, from the public Phi materials.

Audience: Scott Paul, Phi Health / Frequency.

Speak TO Scott as "you" and "your" for the entire episode.

STRICT BANS (the episode fails review if any host breaks these): never refer to Phi Health / Frequency or Scott

in the third person — the words "the/this" + "business/company/owner" and "his/her/their" + "business/company/team"

must not be spoken; say "your business", "your company", or "your operation" instead.

Never mention where this briefing came from, never reference any written input behind the episode,

and never narrate like an outside analyst reviewing a case — you are talking directly with Scott.

Duration target: 9 to 11 minutes. Accept only 8 to 12 minutes. Concise walkthrough, not a lecture.

CTA: complete the qualifier only; no calendar booking language.

12 SECTION AUDIO PACKAGE

1. About Scott and Phi Health / Frequency

The public footprint shows more than a podcast: Phi has an AI health companion, Stack, shop/support paths, research and education content, and now Frequency as a trust-building media layer. The opportunity is to make each expert episode flow into app education, Stack-safe learning, social clips, giveaway follow-up, and support routing without losing scientific caution. A practical AI roadmap for Phi Health: connect Frequency, Stack, the AI Health Journal, research content, giveaway engagement, and support into one repeatable episode-to-education system without drifting into medical overclaims.

2. Current Tool Stack

Phi Health website — Podcast home, product education hub, research entry point, and brand surface. Phi AI Health Journal — Personalized AI health companion and logical destination for episode learning. Frequency podcast — Long-form expert interviews and listener trust engine. Episode transcripts — Raw material for topic maps, clips, app snippets, FAQs, newsletters, and shop-safe education. Phi Stack/shop — Education-to-commerce path that needs careful claims language. Giveaway, support, and app channels — Routing point for listener takeaways, app issues, product interest, partnerships, and medical-advice requests.

3. Industry Context

Phi Health / Frequency operates in consumer AI health, healthspan education, supplement commerce, longevity media, and wellness community. Confirmed audience signals: people pursuing more healthy years, AI-guided wellness app users, Stack customers, and listeners interested in body, mind, recovery, sleep, energy, and expert-led healthspan education. Inferred expansion audience: founders, athletes, action-sports communities, longevity-minded professionals, and health-conscious consumers who want practical guidance without medical overclaiming.

4. Where Time Is Leaking

Expert episodes need repeatable post-production: Your first Frequency episode covers recovery, aging, nutrition, PRP, fascia, muscle, D3/K2, absorption, sleep, and a live Stack giveaway. That is valuable, but the follow-up assets need to be structured carefully so they remain educational. Health claims need a review pass before reuse: A phrase that works in a conversation can become too strong when clipped into an ad, email, shop page, or app note. A review agent flags treatment, supplement, diagnosis, dosing, and outcome-adjacent claims before publishing. Listener questions need routing while intent is fresh: A listener may ask about an episode, a guest, the 3-month Stack giveaway, the app, a product, or a personal medical situation. The faster you route the message safely, the more useful the brand feels. Your education can become an app/product learning library: Every episode can feed searchable FAQs, topic clusters, guest pages, app learning prompts, support macros, and Stack-safe education if it is structured at the transcript level.

5. Six Recommended Agents

Guest Research and Interview Prep Agent, Episode Knowledge Graph Agent, Claims and Citation Review Agent, Clip and Social Repurposing Agent, Listener Intake and Giveaway Ops Agent, App and Stack Learning Path Agent.

6. Agent 1: Guest Research and Interview Prep Agent

Before each Frequency recording, this agent prepares a guest dossier, topic map, interview angles, careful claim notes, and listener-first questions. It helps Scott and the host ask sharper questions without turning the conversation into a generic wellness interview. Outcome for Phi Health / Frequency: Each episode starts with a stronger research base and a cleaner arc.

7. Agent 2: Episode Knowledge Graph Agent

After each episode, this agent tags the transcript by topic, quote, claim risk, evidence need, product relevance, app relevance, and likely listener intent. Outcome for Phi Health / Frequency: A 50-minute episode becomes structured knowledge that can feed content, app education, support, and Stack-safe learning.

8. Agent 3: Claims and Citation Review Agent

This agent reviews episode notes, captions, emails, product education, app content, and giveaway follow-up for claims that need softer language, a citation, a disclaimer, or human review. Outcome for Phi Health / Frequency: Phi can publish more useful education while reducing avoidable claims risk.

9. Agent 4: Clip and Social Repurposing Agent

This agent finds high-value moments in every episode, writes hooks and captions, adapts the same idea for each channel, and keeps language aligned with claims review before posting. Outcome for Phi Health / Frequency: Frequency gets more reach from each recording without making Scott manually cut and rewrite every asset.

10. Agent 5: Listener Intake and Giveaway Ops Agent

When someone comments, replies, enters the Stack giveaway, asks about a guest, or needs the right Phi path, this agent classifies the message and drafts a safe response for review. Outcome for Phi Health / Frequency: Episode engagement gets routed while the listener still cares.

11. Agent 6: App and Stack Learning Path Agent

This agent turns podcast interest into approved next steps across Phi app prompts, Stack education, FAQs, shop support, or human escalation. Outcome for Phi Health / Frequency: Listeners get useful next steps while Phi keeps product and health language inside approved educational boundaries.

12. ROI Picture and Next Step

Use Scott's stated numbers only: 100 monthly leads, \$120 average value, 5% close rate, about 20 hrs weekly admin hours. Conservative language, no guarantees. Close by inviting Scott to open the written playbook and complete the qualifier. The written playbook qualifier lives at <https://bennett-maxwell.github.io/fki-preview/qualify.html> — mention it as the qualifier page, never spell the URL aloud. No calendar-booking language.